

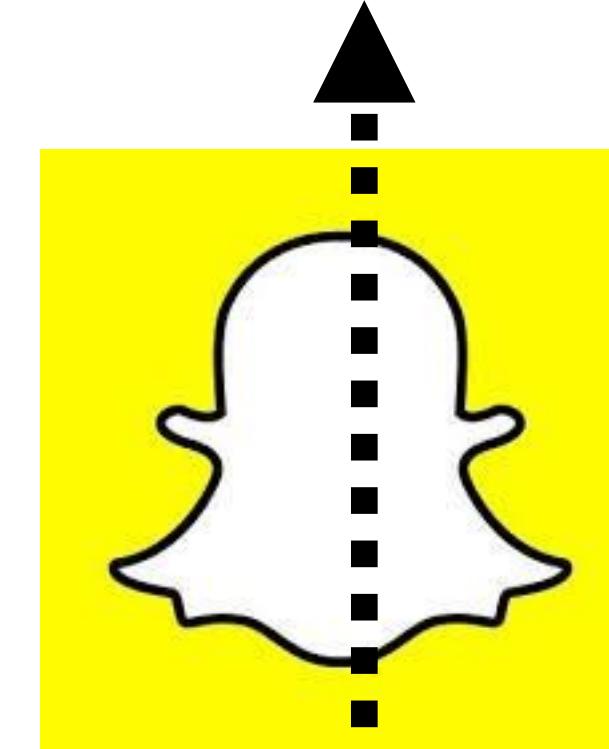
Intro to Databases

SQL

What is a database?

Things that hold info

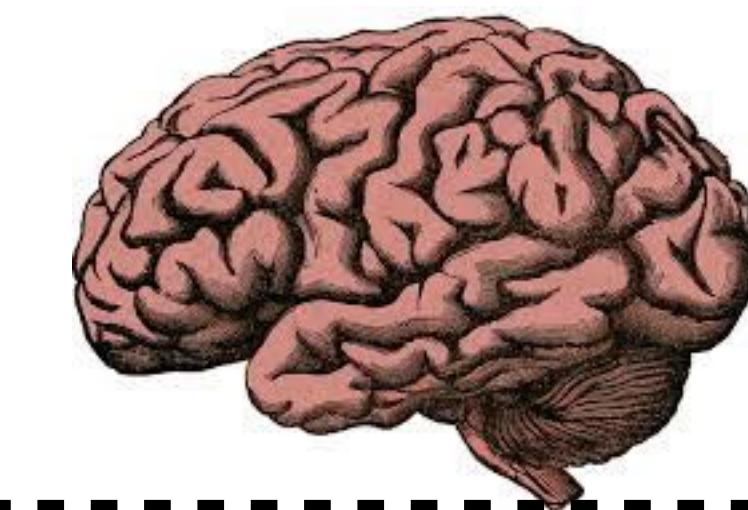
Accessible



Persistent



Product Launch						
	Area	Item Price	Orders	Gross Income	Discount	Net Income
1	Central	\$40	1,800	\$64,000	10	\$6,400.00
2	London	\$40	1,200	\$48,000	7	\$3,360.00
3	North East	\$40	1,200	\$48,000	5	\$2,400.00
4	North West	\$40	1,100	\$44,000	5	\$2,200.00
5	South East	\$40	988	\$39,520	10	\$3,952.00
6	South West	\$40	1,010	\$40,400	5	\$2,020.00
7						\$38,380.00
8			Minimum Gross Income		\$39,520	
9			Maximum Gross Income		\$64,000	
10			Total Net Income		\$217,988	
11			Average Net Income		\$43,596	
12						



A database **persists** information
and is **accessible** via code

organized

queryable

manageable

Organized: Standard Storage Formatting

- DBs are a collection of Tables (or *relations*)
- Tables have Columns (*attributes / fields*) that describe Rows (*instances / tuples*)
- Duplicate rows are not allowed
- Rows often have a primary key (unique identifier)

Table / Relation

Column / Attribute / Field	Column / Attribute / Field	Column / Attribute / Field	
ID	Name	Type	
Row / Tuple / Instance	1	Pikachu	lightning
Row / Tuple / Instance	2	Squirtle	water
Row / Tuple / Instance	3	Charmander	fire
Row / Tuple / Instance	4	Bulbasaur	grass

Queryable: via a Standard Language

- A simple, structured query language: SQL
 - Declarative (vs. imperative)
 - No more hand-rolled algorithms / data structures
 - DBMS picks an efficient execution strategy based on indexes, data, workload etc.



SQL

```
-- Pikachu, I choose you!
SELECT id, name
  FROM pokemon
 WHERE type = 'lightning'
   LIMIT 1
```

Manageable: Easy, Safe, Performant

- Offloads work and requisite understanding of programming
- Knowledge is portable
- Abstraction
- Transfer data between systems
- DBMS can make certain guarantees
 - prevent unsafe operations
 - built-in redundancies
 - handle multiple users, threads

ACID Guarantees

- **Atomicity**
- **Consistency**
- **Isolation**
- **Durability**

Atomic Transactions

- **atomic transaction:** A set of database operations that must occur together
 - i.e. A debit to one bank account, and a credit to another
- A transaction must either succeed or fail; it cannot partially complete.
- Every database query is represented by a transaction

Consistency

- **Specify rules that columns need to follow**
 - Gender column can only contain M, F, or U.
 - Savings account must start with S or checking with C
 - Column cannot be null
- **Protect the database from inconsistencies and simplify software logic**
 - Allows software to make assumptions about underlying data

Resource Management

- ◎ Processes can be readers and writers
- ◎ Files can have many readers
- ◎ If a process has a writer, no other process can read from it, and no other process can write to it

Proposed File Scheme

- Suppose that we have decided not to use a database and instead store our data in a series of files.
- How might our setup fail to serve queries from multiple users?

Deadlock



Databases give us concurrency (Isolation)

- Multiple clients can make queries to read and update without the risk of deadlock or starvation.

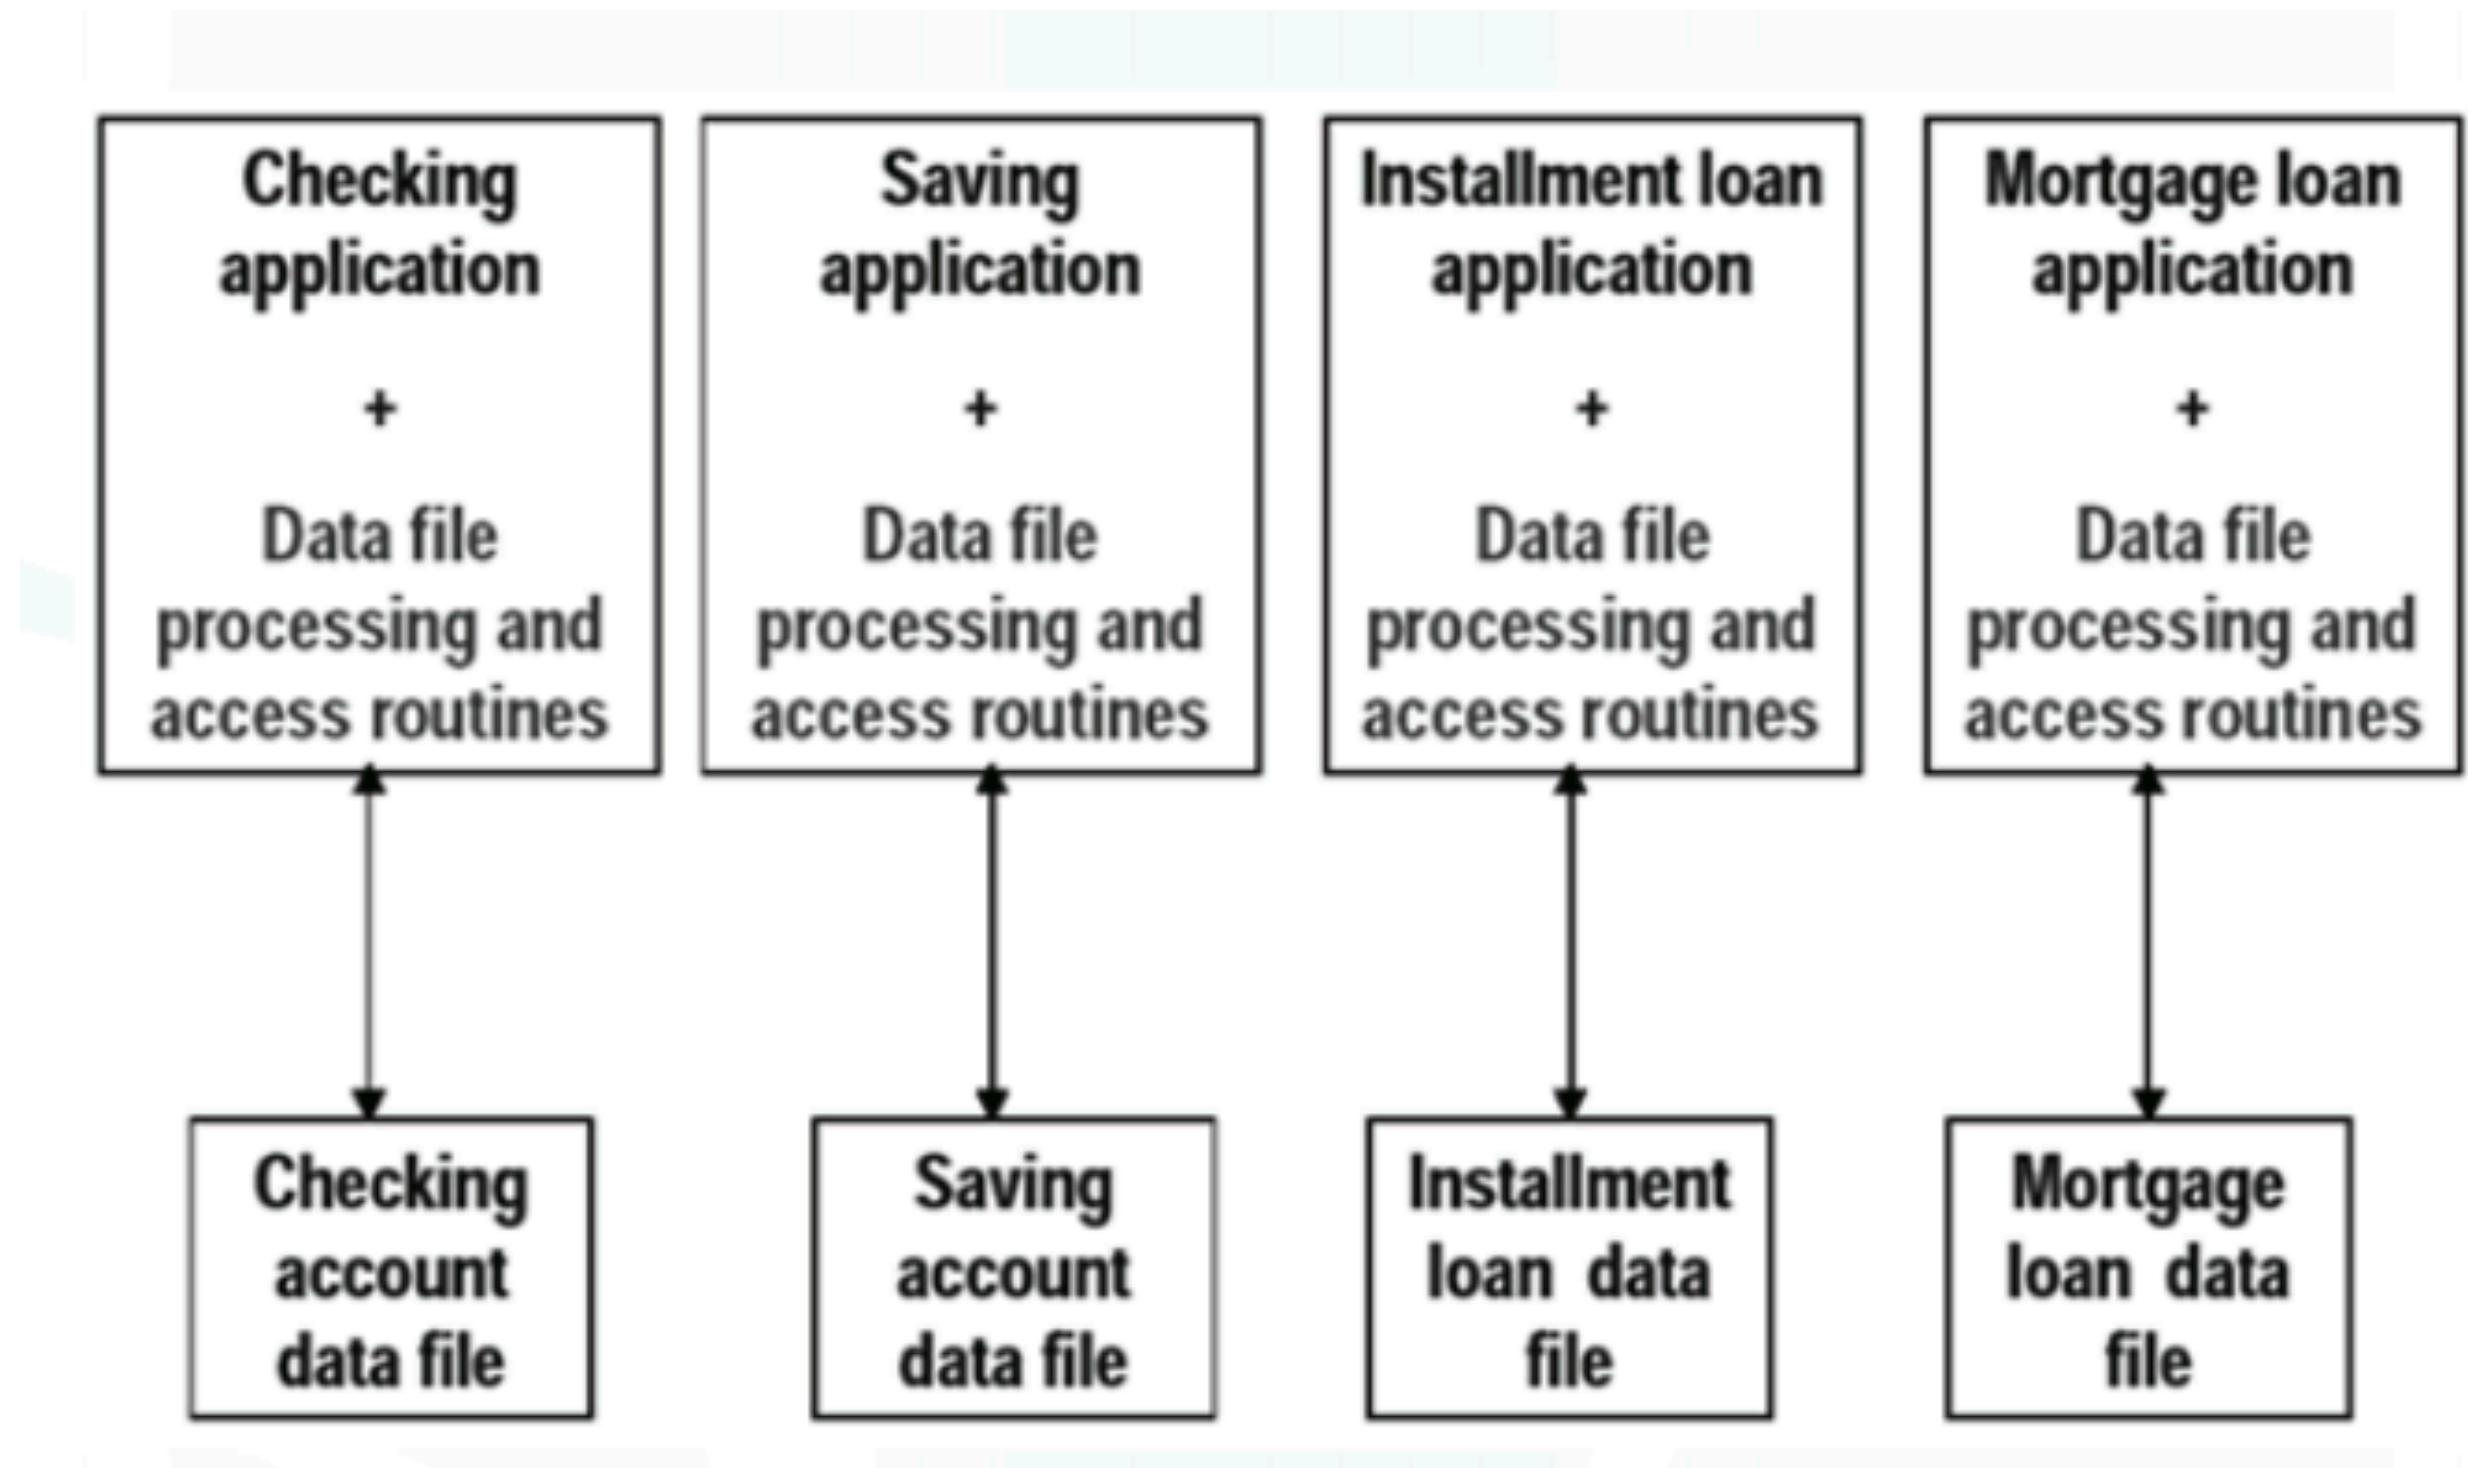
Persistence/Durability

- ◎ Files are also persistence (store information without power)

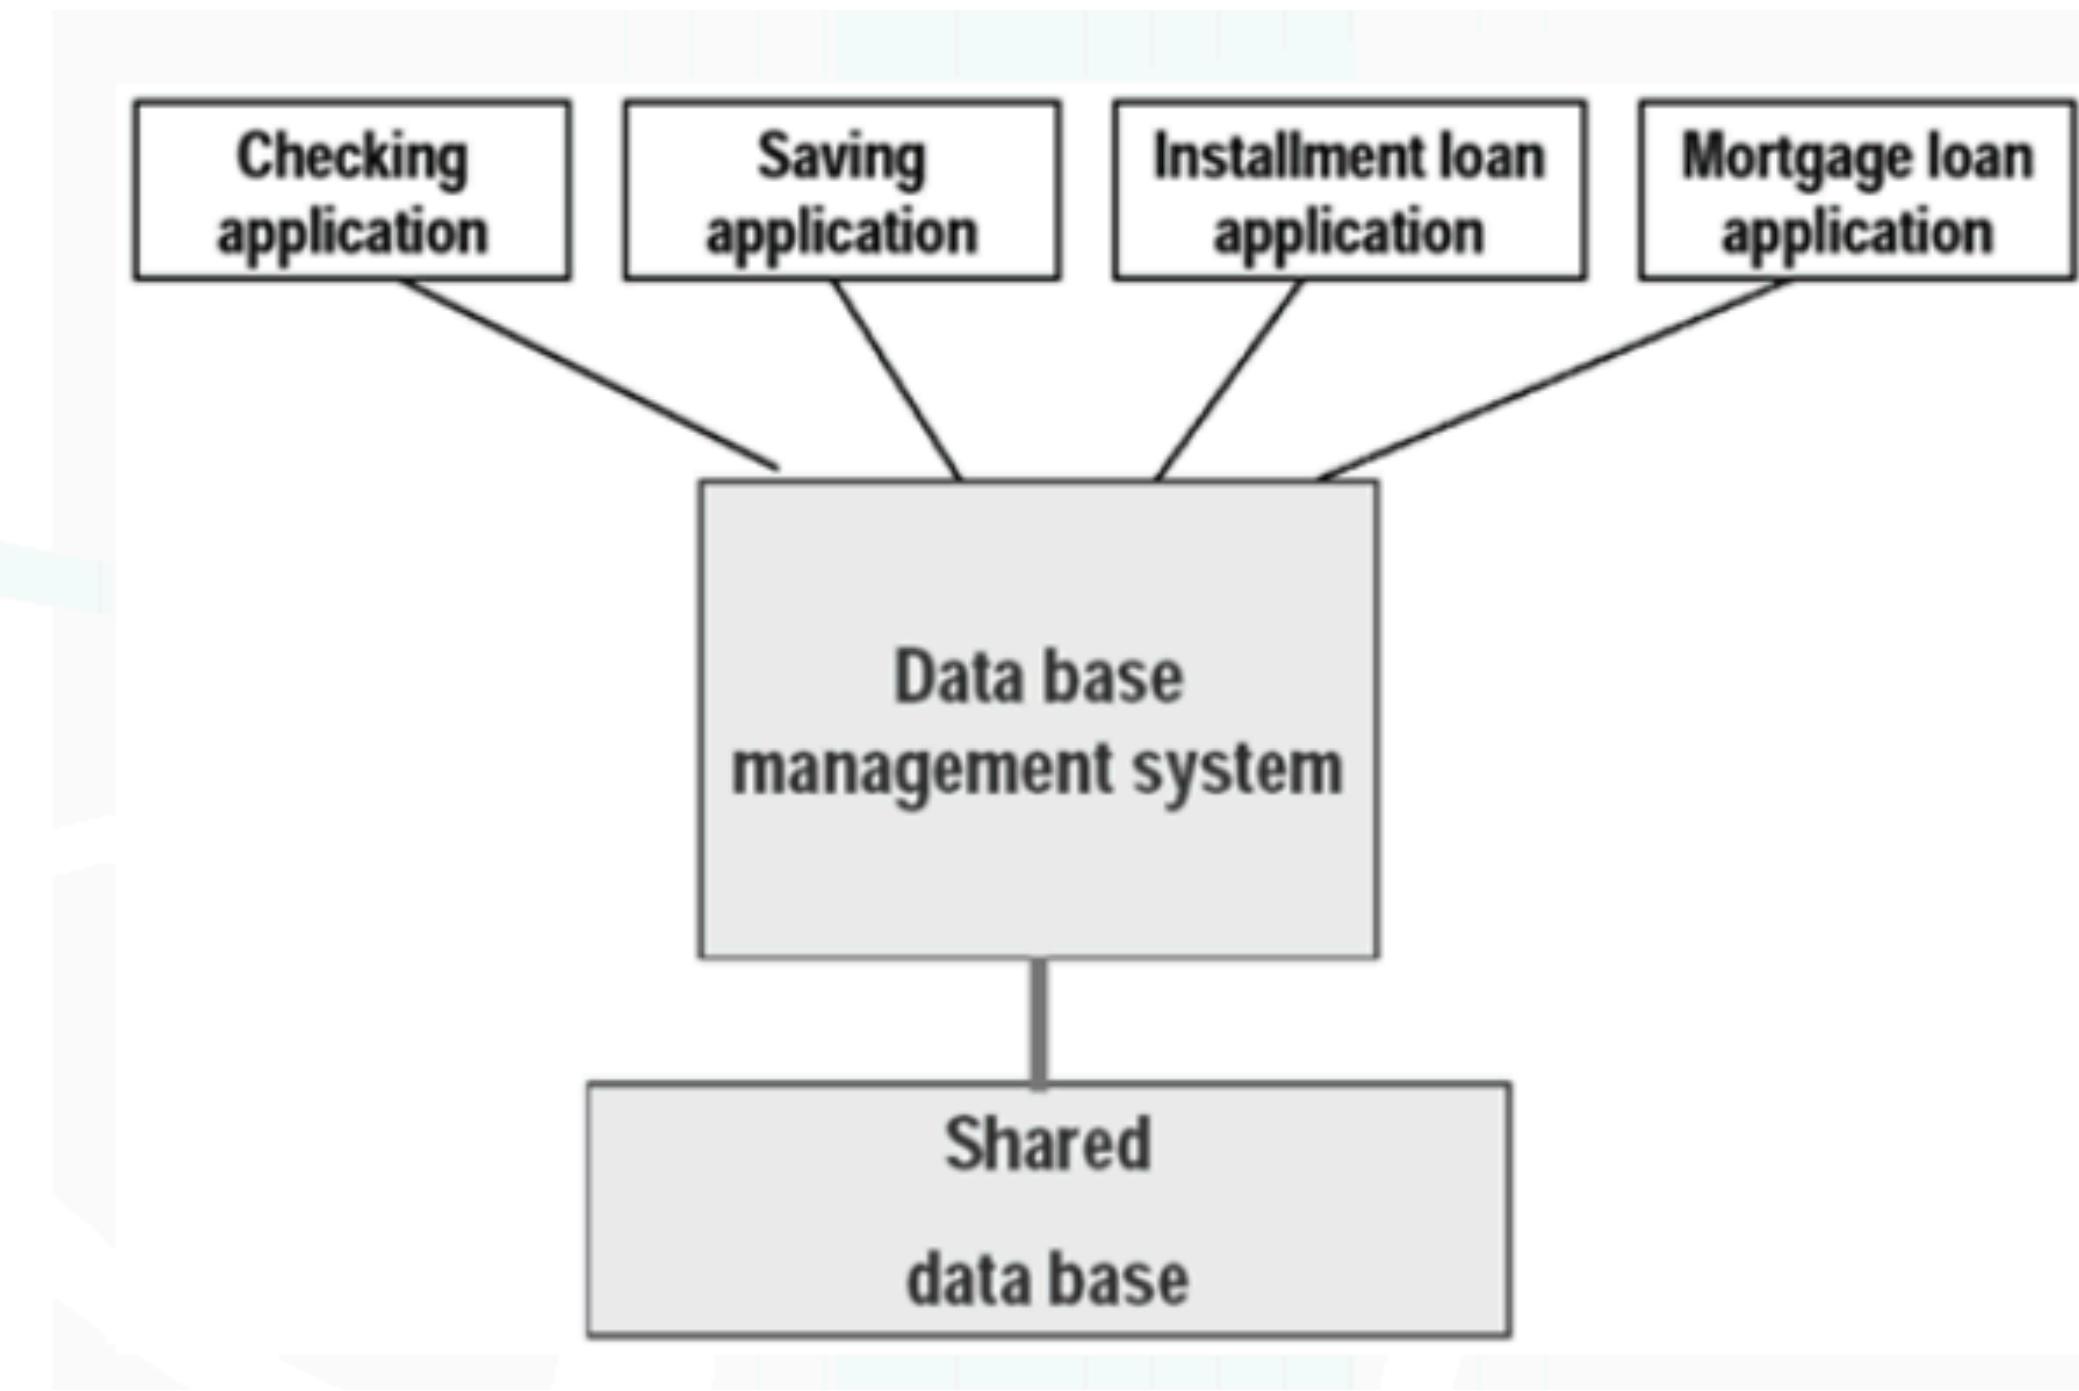
How Did We End Up Here?

Before Relational DBs (ca. < 1970s)

- Data stored in custom “data files”
- Queried via application-specific code
- Advantages
 - Middle layer not needed
 - Solutions customized for each application
- Disadvantages
 - Hard to change the system
 - Knowledge not compounding
 - Data-transfer is difficult



Database Management Systems (DBMS)



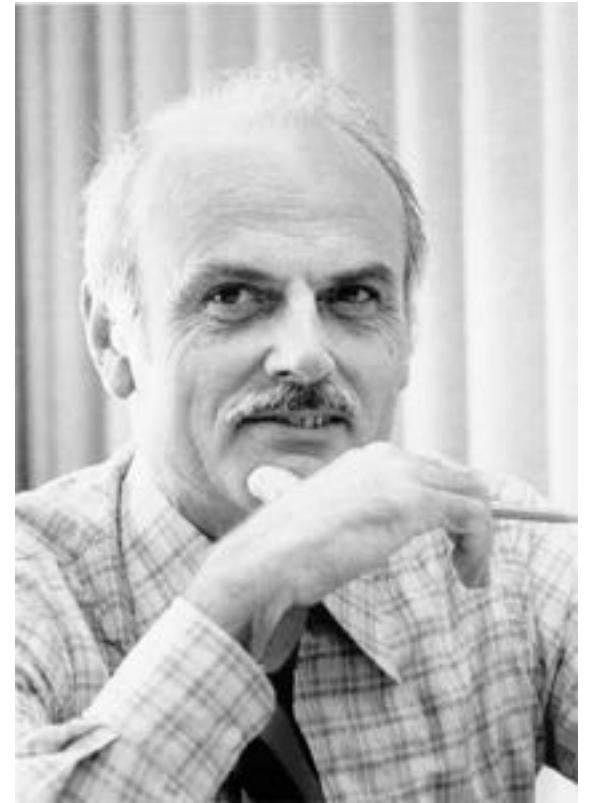
- One layer and language to store and access data
- Sold as a way for “non-technical people” to manage data

“Future users of large data banks must be protected from having to know how the data is organized in the machine (the internal representation).”

– E. F. CODD,

A RELATIONAL MODEL OF DATA FOR
LARGE SHARED DATA BANKS

Relational Databases & Logic



- 1969: Edgar Frank "Ted" Codd outlines *relational model* of data
- Wrote Alpha (never implemented) as a *query language*
- IBM slow to adopt his ideas
 - Competitors started to do so
 - IBM team formed without Codd, created **Structured English Query Lang**
- SEQUEL way better than what came before
 - 1979: copied by Larry Ellison (from pre-launch papers / talks!) as "SQL"
- SQL became the standard (ANSI 1986, ISO 1987)
 - Codd continued to fault SQL compared to his theoretical model
 - The Third Manifesto: solve the *object-relational impedance mismatch*

Appreciating Databases

- **Ubiquitous**
- **Standardized**
- **Complex / deep**
- **Powerful: database admins are**
 - Feared by developers
 - ...but also taken for granted until things break
 - Befriended by business people
 - Contacted by the government for secret data (e.g. NSA)

Progression of Databases

- **Navigational (< 1970s)**
 - More common during tape era; entries had references to next entries.
- **Relational (> 1970s)**
 - Based on relational (table-based) logic, see E.F. Codd.
- **NoSQL (> 2000s)**
 - "Not only SQL" — document storage, for example.

RDBMS vs NoSQL

- **A DBMS doesn't have to be relational**
 - Remember, DBMS is just an application that intelligently stores data and can answer requests to manage that data
- **Lately, many "NoSQL" or non-relational DBMSs have been gaining popularity**
 - Graph databases (e.g. Neo4J)
 - Document databases (e.g. MongoDB)
 - Hybrids (e.g. PostgreSQL)
- **RDBMSs still remain the #1 DB option for now**



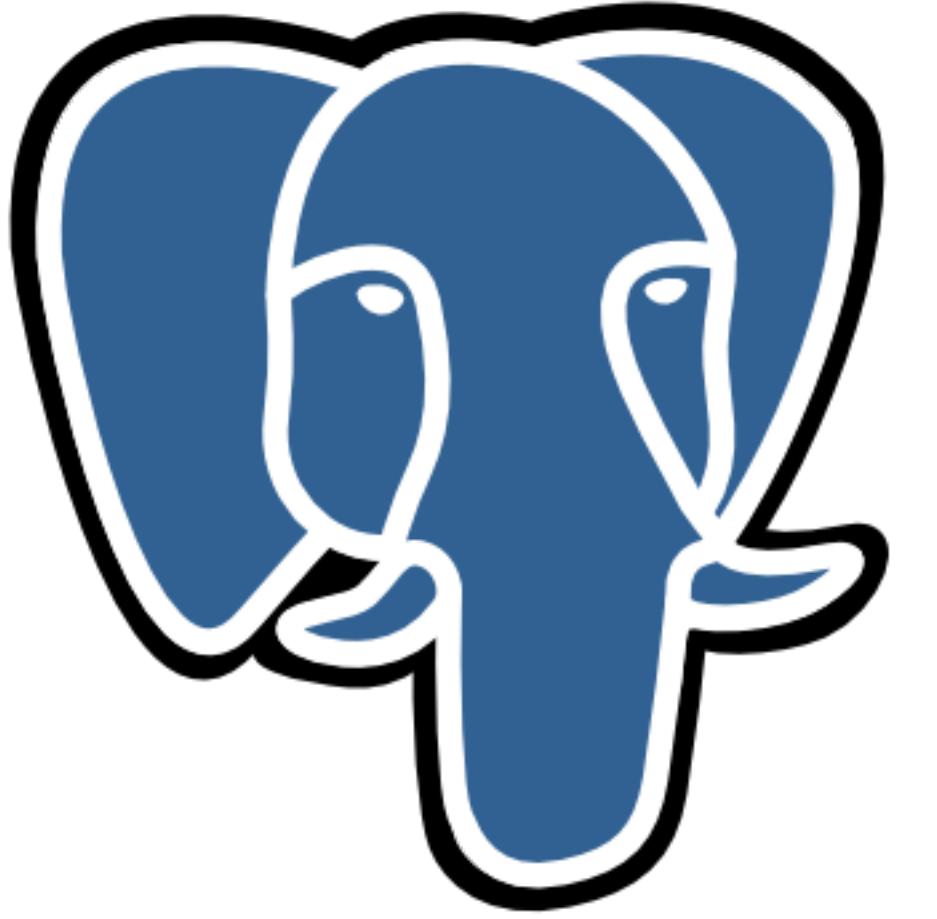
Some well-known rDBMSs



PostgreSQL



MySQL®



Postgre**SQL**

Why PostgreSQL?

- Advanced, powerful, and popular
- Rapid open source development
- Highly extensible (stored procedures)
- Deep SQL standards compliance
- NoSQL ("Not Only SQL"), objective support
- Excellent transactions / ACID reliability; focus on integrity
- Multi-user management / administration

History of PostgreSQL

- 1970s at UC Berkeley:
INteractive Graphics REtrieval System (INGRES)
- 1980s: POSTGRES ("Post-Ingres")
- 1995: POSTQUEL and Postgres95.
 - monitor -> psql
- 1996: Adopted by the open source community
 - Ongoing: stability, testing, documentation, new features
 - PostgreSQL

psql

```
psql [local]:5432 glebec # \l
List of databases
+-----+-----+-----+-----+-----+-----+
| Name | Owner | Encoding | Collate | Ctype | Access privileges |
+-----+-----+-----+-----+-----+-----+
| assessmentexpresssequelize | fullstack | UTF8 | en_US.UTF-8 | en_US.UTF-8 | 
| auther | glebec | UTF8 | en_US.UTF-8 | en_US.UTF-8 | 
| checkpoint_angular | glebec | UTF8 | en_US.UTF-8 | en_US.UTF-8 | 
| checkpoint_express_review | glebec | UTF8 | en_US.UTF-8 | en_US.UTF-8 | 
| glebec | glebec | UTF8 | en_US.UTF-8 | en_US.UTF-8 | 
| juke | glebec | UTF8 | en_US.UTF-8 | en_US.UTF-8 | 
| postgres | glebec | UTF8 | en_US.UTF-8 | en_US.UTF-8 | 
| sequelizecheckpoint | fullstack | UTF8 | en_US.UTF-8 | en_US.UTF-8 | 
| template0 | glebec | UTF8 | en_US.UTF-8 | en_US.UTF-8 | =c/glebec
| template1 | glebec | UTF8 | en_US.UTF-8 | en_US.UTF-8 | =c/glebec
| triplanner | glebec | UTF8 | en_US.UTF-8 | en_US.UTF-8 | =c/glebec
| twitterdb | glebec | UTF8 | en_US.UTF-8 | en_US.UTF-8 | 
| wikistack | glebec | UTF8 | en_US.UTF-8 | en_US.UTF-8 | 
+-----+-----+-----+-----+-----+-----+
(13 rows)

psql [local]:5432 glebec # \c auther
You are now connected to database "auther" as user "glebec".

psql [local]:5432 glebec # auther # \dt
List of relations
+-----+-----+-----+-----+
| Schema | Name | Type | Owner |
+-----+-----+-----+-----+
| public | stories | table | glebec |
| public | users | table | glebec |
+-----+-----+-----+-----+
(2 rows)
```



pgcli

```
stayupdated_test> \d
+-----+-----+-----+-----+
| Schema | Name          | Type   | Owner  |
+-----+-----+-----+-----+
| public  | admins         | table   | amjith |
| public  | cpes           | table   | amjith |
| public  | goose_db_version | table   | amjith |
| public  | goose_db_version_id_seq | sequence | amjith |
| public  | packages        | table   | amjith |
| public  | packages_id_seq | sequence | amjith |
| public  | users           | table   | amjith |
| public  | users_id_seq    | sequence | amjith |
| public  | vulnerabilities  | table   | amjith |
| public  | vulnerabilities_cpes | table   | amjith |
| public  | vulnerabilities_id_seq | sequence | amjith |
+-----+-----+-----+-----+
SELECT 11
stayupdated_test> SELECT * FROM users;
+-----+-----+-----+-----+
| id | display_name | password | email           | created_on      |
+-----+-----+-----+-----+
| 177 | DisplayName1 | 1024cms  | user@ex.com    | 2014-11-15 15:02:50.094560 |
| 180 | testname2    | pas5w0rd | email@ex.com   | 2014-11-28 10:25:46.170660  |
| 181 | amjith       | password  | amjith@amjith.amjith | 2014-11-28 18:39:48.195067 |
+-----+-----+-----+-----+
SELECT 3
stayupdated_test> SELECT * FROM [REDACTED]
[REDACTED]
admins
cpes
goose_db_version
packages
users
```

Postico

The screenshot shows the Postico PostgreSQL client interface. The main window displays a table of data from the 'forex' table in the 'reporting' database. The table has columns: currency, base, rate, date, and source_id. The data includes various currencies like AED, AFN, ALL, AMD, ANG, AOA, ARS, AUD, AWG, AZN, BAM, BBD, BDT, and BGN, all converted to USD. The 'date' column shows 2014-07-25 for most entries. The 'source_id' column shows values like 1, 2, and 3. On the left, a sidebar lists other tables in the schema: SQL Query, currencies, daily_proceeds, forex, forex_sources, fs_daily_proceeds, fs_daily_sales, fs_orders, fs_proceeds_weekly, fs_weekly_sales, itc_daily_proceeds, itc_daily_sales, itc_proceeds_weekly, itc_reports, itc_reports_daily, itc_reports_monthly, and itc_reports_weekly. The 'forex' table is currently selected. At the bottom, there are buttons for Content, Structure, Filter, and navigation.

Reporting > reporting > forex Connected. PostgreSQL 9.4.5

	currency	base	rate	date	source_id	
currencies	AED	USD	3.67291	2014-07-25	1	<input type="button" value="grid"/>
daily_proceeds	AFN	USD	56.485726	2014-07-25	1	<input type="button" value="grid"/>
forex	ALL	USD	103.5838	2014-07-25	1	<input type="button" value="grid"/>
forex_sources	AMD	USD	410.086	2014-07-25	1	<input type="button" value="grid"/>
fs_daily_proceeds	ANG	USD	1.787	2014-07-25	1	<input type="button" value="grid"/>
fs_daily_sales	AOA	USD	96.952626	2014-07-25	1	<input type="button" value="grid"/>
fs_orders	ARS	USD	8.169642	2014-07-25	1	<input type="button" value="grid"/>
fs_proceeds_weekly	AUD	USD	1.062844	2014-07-25	1	<input type="button" value="grid"/>
fs_weekly_sales	AWG	USD	1.79	2014-07-25	1	<input type="button" value="grid"/>
itc_daily_proceeds	AZN	USD	0.784067	2014-07-25	1	<input type="button" value="grid"/>
itc_daily_sales	BAM	USD	1.453024	2014-07-25	1	<input type="button" value="grid"/>
itc_proceeds_weekly	BBD	USD	2	2014-07-25	1	<input type="button" value="grid"/>
itc_reports	BDT	USD	77.61971	2014-07-25	1	<input type="button" value="grid"/>
itc_reports_daily	BGN	USD	1.452543	2014-07-25	1	<input type="button" value="grid"/>
itc_reports_monthly						
itc_reports_weekly						

Content Structure + Filter < Page 1 of 342 >

SQL Query currency base rate date source_id

SQL Query currency
MULTIPLE

SQL Query base
USD

SQL Query rate
MULTIPLE

SQL Query date
2014-07-25

SQL Query source_id forex_source:
id name
1 Open Exchange R...

LAB